

Sharan Chenna

Site Reliability Engineer

me@sharanch.dev · [LinkedIn](#) · [GitHub](#) · +91 82976 22321 · Hyderabad, India

SUMMARY

SRE with 4+ years building and operating large-scale Linux and OCI Compute infrastructure. First hire on Oracle OCI's CloudOps SRE team, automated 15+ manual runbooks via a Python CLI, cutting resolution time from 45+ minutes to under 10. Ramped a team of 6 and personally owned 24x7 P0/P1 response. Currently on a deliberate sabbatical shipping production-grade open source tooling in AIOps, Kubernetes observability, and cloud-native reliability engineering.

SKILLS

Infrastructure: OCI, AWS, Kubernetes, Terraform, Linux, Docker, Helm, ArgoCD, KVM

Observability: Prometheus, Grafana, Alertmanager, Loki, Zabbix

Automation/CICD: Python, Bash, Go, GitHub Actions, Jenkins, Ansible, Boto3, Terraform

Reliability: Incident command (P0/P1), post-mortems, SLO design, chaos engineering, on-call operations

EXPERIENCE

Cloud Operations Engineer · Oracle · OCI Compute

Feb 2024 – Sep 2025 · Hyderabad

- **Defined and enforced SLOs, error budgets, and alerting thresholds** for OCI Compute Control and Data Plane, reducing high-noise alert volume from 40+ per week to 3–5 through redesigned Prometheus alerting rules.
- **Led incident response, post-mortems, and systemic improvements** to reduce MTTR by automating 15+ of ~40 manual runbooks via a Python CLI, cutting resolution time from 45+ minutes to under 10.
- **Built and maintained internal tooling** for runbook automation, self-healing workflows, and capacity management scripts using Python, Ansible, and Terraform.
- **Designed and operated Kubernetes-based infrastructure** at scale across multiple regions, implementing Terraform taint/replace for automated fleet recovery from hardware failures.
- **Instrumented services with Prometheus metrics, structured logging (Loki), and distributed tracing** for observability, creating Grafana dashboards used as the team's primary incident response interface.
- **Managed on-call rotations, triaged production alerts, and owned escalation paths** for thousands of bare-metal and VM nodes, leading triage and resolution of P0/P1 outages.

Associate Engineer · CtrlS Datacenters · Infrastructure Engineering

Apr 2021 – Dec 2023 · Hyderabad

- **Defined and enforced SLOs, error budgets, and alerting thresholds** for 90+ enterprise private clouds, maintaining 99.99% uptime SLA through proactive monitoring and structured post-incident reviews.
- **Built and maintained internal tooling** for runbook automation, self-healing workflows, and capacity management scripts using Bash, Python, and YUM/DNF for automated patch orchestration.
- **Designed and operated Kubernetes-based infrastructure** at scale, implementing Zabbix for observability with custom alerting rules and per-client dashboards to reduce MTTD for critical failures.
- **Instrumented services with Prometheus metrics, structured logging (Loki), and distributed tracing** for observability, creating custom alerting rules and thresholds to enable faster cross-team escalation during incidents.
- **Partnered with engineering teams to enforce observability standards** and improve deployment pipelines via GitOps with ArgoCD and multi-stage Alpine builds reducing image size by ~60%.

OPEN SOURCE ENGINEERING

Self-directed · Open Source

Deliberately stepped back from employment to build depth in cloud-native, AIOps, and observability engineering. Shipped production-grade projects independently, each designed to demonstrate real operational patterns — not toy demos.

[log-explainer](#) — [Python](#) · [Ollama](#) · [GitHub Actions](#) · [GHCR](#)

- **Built AIOps CLI for real-time log analysis** using locally running LLMs (Ollama, qwen2.5-coder:1.5b) with two-pass severity classification and incident summarization on pattern anomalies.
- **Instrumented full observability stack** with Prometheus metrics, structured logging, and alerting pipelines for end-to-end monitoring and SLO tracking.

[go-sre-observatory](#) — [Go](#) · [Kubernetes](#) · [Prometheus](#) · [Grafana](#) · [Loki](#)

- **Designed and operated Kubernetes-based infrastructure** at scale with full RED metrics instrumentation, SLI definition, and SLO-breach simulation to keep alerting pipelines continuously exercised.

- **Built and maintained internal tooling** for runbook automation, self-healing workflows, and capacity management scripts using Go and Prometheus alerting rules.

[chatops](#) — *React · Node.js · PostgreSQL · ArgoCD · Helm*

- **Partnered with engineering teams to enforce observability standards** and improve deployment pipelines via GitOps with ArgoCD and multi-stage Alpine builds reducing image size by ~60%.

[istio-mesh-demo](#) — *Istio · Kubernetes · FastAPI · Kiali · Grafana*

- **Designed and operated Kubernetes-based infrastructure** at scale with full mTLS encryption via Envoy sidecars across distributed microservices.
- **Built and maintained internal tooling** for runbook automation, self-healing workflows, and capacity management scripts using Istio for service mesh and canary pipeline shifting.

[postgresql-ha-lab](#) — *CloudNativePG · Kubernetes*

- **Instrumented full observability stack** with kube-prometheus-stack, defining SLIs for replication lag and failover time validated against SLO targets under simulated failure conditions.

EDUCATION & CERTIFICATIONS

Bachelor of Computer Applications · Kakatiya University

2020 · 8.23 CGPA

- **Active:** OCI Foundations Associate · OCI AI Foundations · LFS162 DevOps & SRE · GitHub Professional Certificate
- **In progress:** AWS Solutions Architect – Associate (SAA-C03)